

Lecture 10 – Variational Autoencoders

Provide short answers (max. 25 words each) to the following questions.

- What is the goal of latent variable models?
- What does marginalization in latent variable models achieve?
- What is the prior distribution typically used for continuous latent variables in VAEs?
- What is Ancestral Sampling in generative models?
- What is the main challenge in training Latent Variable Models?
- What is ELBO and why is it used?
- What are the two main terms in ELBO?
- What does the KL divergence term enforce?
- What does the encoder in a VAE output?
- What does the decoder in a VAE do?
- Why is Monte Carlo estimation used in VAEs?
- What is the total VAE loss function?
- What does reconstruction loss measure?
- What does KL divergence measure in VAEs?
- How can VAEs generate new samples?
- What is disentanglement in VAEs?
- What is a Regularized Discriminative VAE?
- What additional term is added in Regularized Discriminative VAE loss?
- What does repulsive loss enforce?
- What are the components of the total loss in Regularized Discriminative VAE?
- How do VAEs differ from standard Autoencoders?
- Why are VAEs considered generative models?

Multiple-choice questions

What is the goal of Latent Variable Models?

- A) Predict labels directly
- B) Learn a probability distribution over observed data using hidden variables
- C) Remove noise from data
- D) Perform clustering only

What does marginalization in Latent Variable Models achieve?

- A) Computes latent variables directly
- B) Removes prior distribution
- C) Computes data likelihood by integrating over latent variables
- D) Normalizes embeddings

Which of the following is an example of a discrete Latent Variable Model?

- A) Variational Autoencoder
- B) GAN
- C) Mixture of Gaussians
- D) Normalizing Flow

What prior distribution is typically used for continuous latent variables in VAEs?

- A) Uniform distribution
- B) Standard Normal distribution
- C) Bernoulli distribution
- D) Exponential distribution

What is Ancestral Sampling?

- A) Sampling latent variables from the prior and generating data through the decoder
- B) Sampling noise from a GAN
- C) Sampling from posterior only
- D) Removing latent space

What is the main challenge in training Latent Variable Models?

- A) Lack of latent space
- B) Intractable integral for log-likelihood
- C) Too few parameters
- D) No reconstruction loss

What does ELBO stand for?

- A) Evidence Lower Bound
- B) Expected Latent Bound
- C) Encoder Latent Base Output
- D) Error Loss Bound

What are the two main terms in ELBO?

- A) Reconstruction term and KL divergence term
- B) Cross-entropy and hinge loss
- C) Reconstruction term and Wasserstein distance
- D) KL divergence and GAN loss

What does the KL divergence term enforce?

- A) Similarity between approximate posterior and prior
- B) Similarity between encoder and decoder
- C) Diversity in latent space

D) Removal of latent variables

What does the encoder in a VAE output?

- A) Latent variable only
- B) Mean and variance of approximate posterior
- C) Reconstruction error
- D) GAN discriminator score

What does the decoder in a VAE do?

- A) Adds noise to latent variables
- B) Generates data from latent variables
- C) Computes KL divergence
- D) Normalizes embeddings

Why is Monte Carlo estimation used in VAEs?

- A) To approximate expectations in ELBO
- B) To compute GAN loss
- C) To normalize latent space
- D) To enforce Lipschitz constraint

What is the reparameterization trick?

- A) Express sampling as a deterministic function of parameters and noise
- B) Remove latent space
- C) Use GAN generator
- D) Normalize embeddings

What is the total VAE loss function?

- A) Reconstruction loss only
- B) Negative ELBO = Reconstruction loss + KL divergence
- C) GAN loss + KL divergence
- D) Wasserstein loss

What does reconstruction loss measure?

- A) Similarity between posterior and prior
- B) How well the decoder reconstructs input data
- C) Diversity in latent space
- D) GAN discriminator accuracy

What does KL divergence measure in VAEs?

- A) Difference between approximate posterior and prior
- B) Difference between encoder and decoder

- C) Reconstruction accuracy
- D) GAN loss

How can VAEs generate new samples?

- A) Sample latent variables from prior and decode them
- B) Use GAN discriminator
- C) Remove latent space
- D) Apply clustering

What is resynthesis in VAEs?

- A) Apply GAN loss
- B) Remove KL divergence
- C) Normalize embeddings
- D) Modify latent representation and decode to create new variations

How can VAEs perform attribute manipulation?

- A) By removing the latent space
- B) By using GAN discriminator
- C) By computing the difference vector in latent space and apply it
- D) By applying Dropout

What is disentanglement in VAEs?

- A) Learning latent dimensions that correspond to interpretable properties
- B) Removing latent space
- C) Using GAN loss
- D) Applying random noise

What is a Regularized Discriminative VAE?

- A) GAN variant
- B) VAE variant that learns discriminative representations for classification
- C) Autoencoder without latent space
- D) Normalizing flow

What additional term is added in Regularized Discriminative VAE loss?

- A) GAN loss
- B) Repulsive loss
- C) Wasserstein loss
- D) Dropout penalty

What does repulsive loss enforce?

- A) Minimum distance between latent clusters of different classes

- B) Similarity between posterior and prior
- C) Diversity in generated samples
- D) Removal of latent space

What are the components of total loss in Regularized Discriminative VAE?

- A) Reconstruction + Wasserstein loss
- B) GAN loss + KL divergence
- C) Reconstruction + KL + Repulsive loss
- D) KL divergence only

How do VAEs differ from standard autoencoders?

- A) VAEs remove latent space
- B) VAEs learn probabilistic latent representations
- C) VAEs use GAN loss
- D) VAEs ignore reconstruction

Why are VAEs considered generative models?

- A) They can sample new data points from learned distribution
- B) They use GAN discriminator
- C) They remove latent space
- D) They apply clustering only